

# OpenMP Offloading

-- getting started  
through examples --

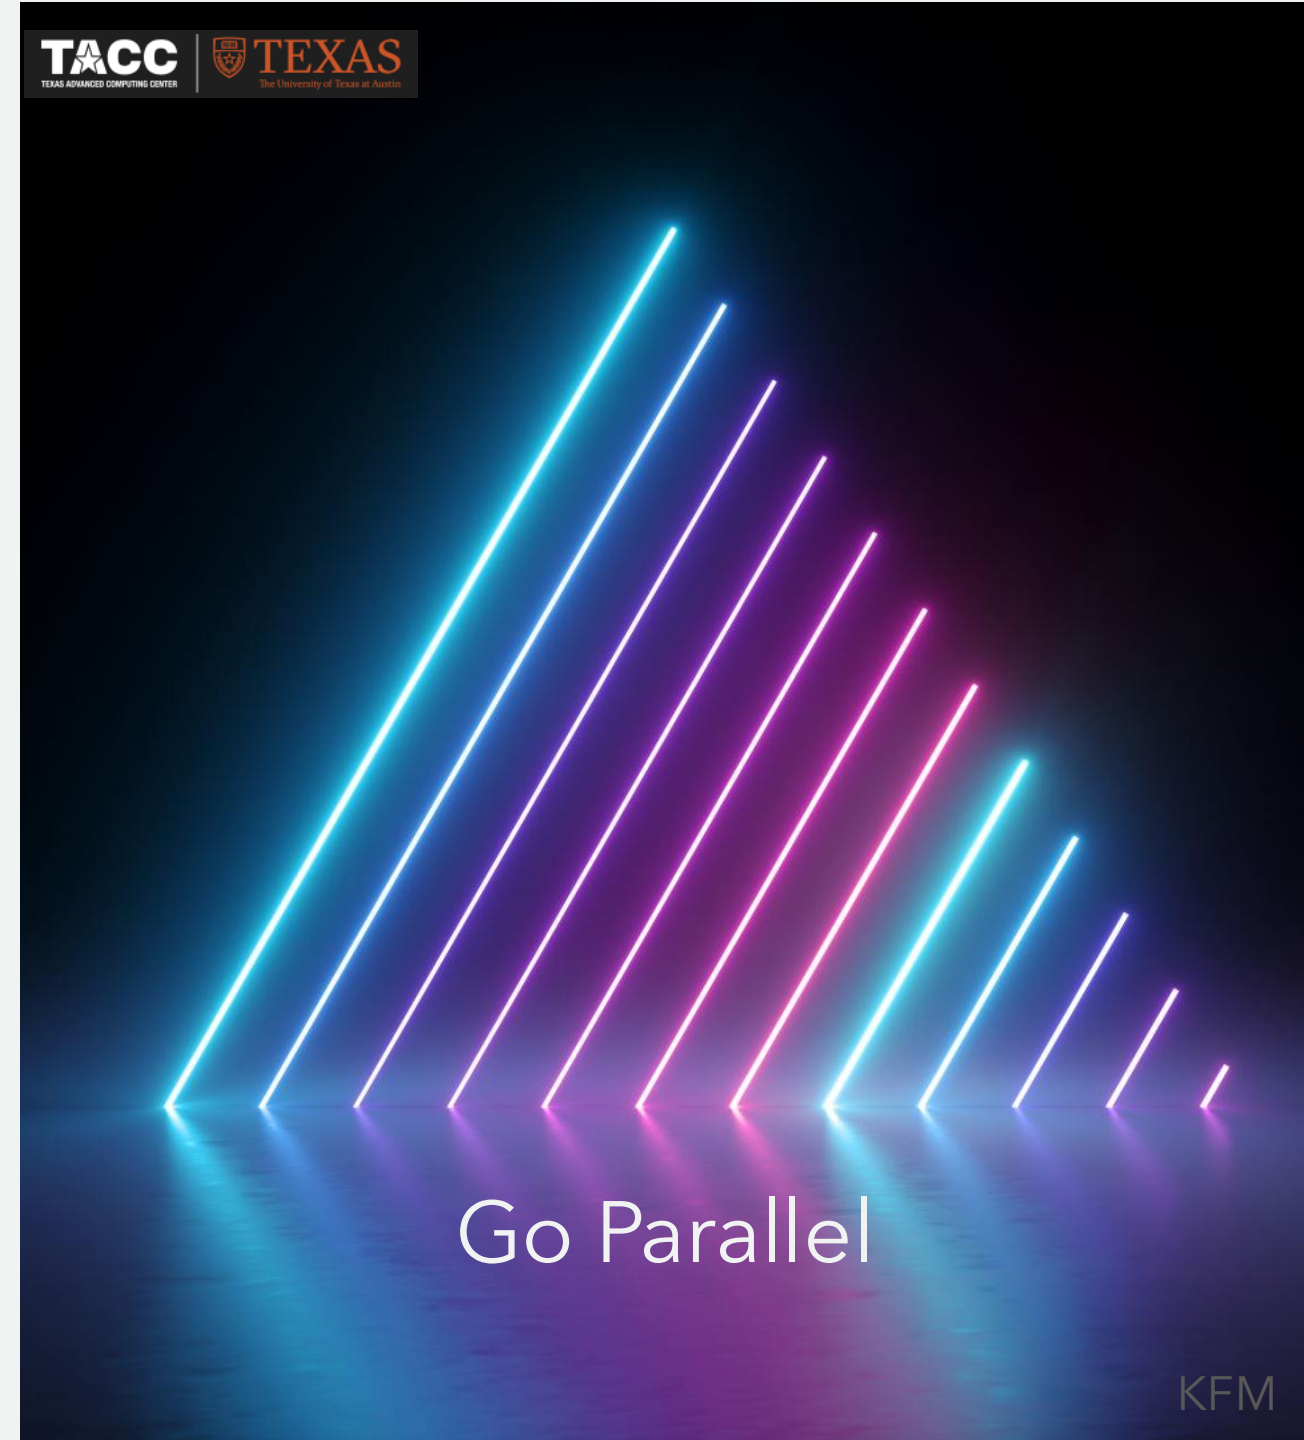
by

Kent Milfeld

milfeld@tacc.utexas.edu

<https://tinyurl.com/tacc-2025-appi>

[https://github.com/milfeld/openmp\\_offload](https://github.com/milfeld/openmp_offload)



Go Parallel



Required Proficiency: :

- Some experience or basic knowledge of OpenMP
- Working knowledge of C/C++ or F90
- Basic GPU Architecture Understanding

References: [openmp.org/specifications](https://openmp.org/specifications)

6.0 Specification

6.0 Examples

Basic operations/concepts of code are explained in tutorial.

(Non-tutorial viewers should read grammar/syntax text after code sections first, then peruse code.)



# Agenda

We will not cover Unified Shared Memory.

## ■ OpenMP

- Background and Levels of Parallelism.
- GPU Programming.
- Review Worksharing Loop.
- Quick Intro to Simple Tasking.

Useful for understanding modern OpenMP, but not necessary for basic offloading.

## OpenMP Offloading

- Launch and direct device execution with **target**.
- Explicit/Implicit **mapping** ("transfer") to/from device.
- **teams**, **distribute**, and **loop**.
- Data Persistence: **target data**.  
Data Movers: **target enter/exit data**, **target update**.
- **declare target** (ensure function/data execution/access on device).  
Asynchronous Offloading with **nowait** and **depend**.

# OpenMP -- board and membership

Since 1997

A non-profit organization

Governed by OpenMP Architecture Review Board (ARB)





# OpenMP

## deals with multiple levels of Parallelism

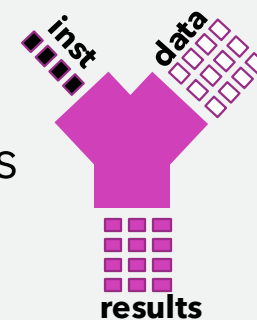
parallel do/for, task

Across Cores



simd

Across SIMD Lanes



target

Across Devices



# Modern OpenMP -- beyond Worksharing Loops

Just a few directives and clauses.

Version	Directives	Parallelism
pre-3.0	parallel for do	worksharing loops
3.0	<b>task</b>	explicit tasking, <b>dependence</b> scheduling
4.0	<b>simd</b>	SIMD
	<b>target</b>	device offloading
	( <b>teams distribute,</b> <b>loop</b> and <i>data-mapping</i> )	
:		
6.0	...	

Other important directives and clauses.

CPU affinity (5.0)  
loop transforms (5.1-6.0)  
free-agent threads (6.0)  
taskgraphs (6.0)

# Computing on GPU (with separate host and device storage)

h\_ == host   d\_ == device

```
declare foo a gpu cuda kernel
```

```
main() {
```

```
declare h_x,   d_x (pointer)
```

```
...
```

```
    alloc_on_dev( sz            d_x )
```

```
    copy_to_dev(    h_x -> d_x )
```

```
    foo(d_x)
```

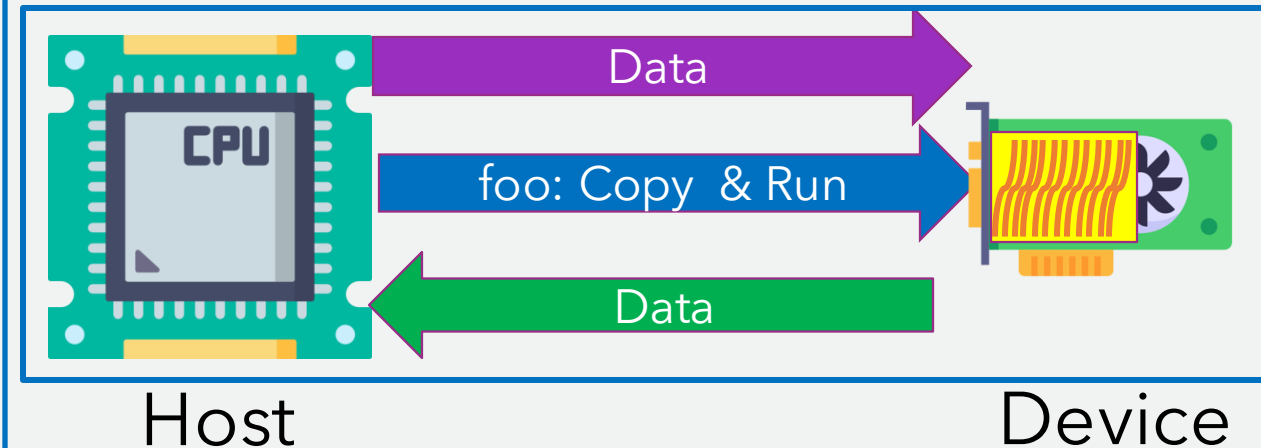
```
    copy_from_dev( h_x <- d_x )
```

```
    dealloc_on_dev(            d_x )
```

```
...
```

**compile prog.c** # creates fat binary

**./a.out**        # run on host, exec foo binary on gpu



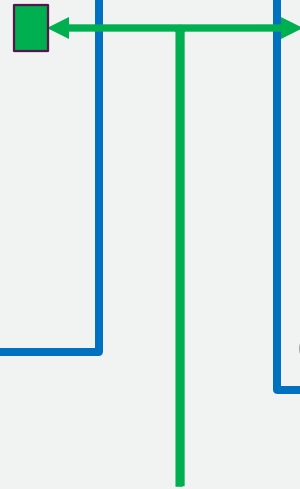
CUDA programs look like this, where it all began.



# Worksharing Loop

```
int main(){  
    float a=9.0, x[16];  
  
    #pragma omp parallel for  
    for(int i=0; i<16; i++)  
        x[i]=sqrt(a*i*i);  
  
}
```

```
program main  
    real :: a=9.0, x(16)  
  
    $omp parallel do  
        do i = 1,16  
            x(i)=sqrt(a*i*i)  
        end do  
  
    end program
```



creates *team* of threads (e.g. export OMP\_NUM\_THREADS=4)

*implicit tasks* (iteration chunks of work) executed by threads



# Worksharing Loop

```
void foo(float a, float *x, int n){  
  
    printf("id=%d\n",omp_get_thread_num());  
  
    #pragma omp for  
    for(int i=0; i<n; i++)  
        x[i]=sqrt(a*i*i);  
}  
  
int main(){  
    float a=9.0, x[16];  
    #pragma omp parallel  
    {  
        foo(a,x,16);  
    }  
}
```

```
subroutine foo(a, x, n)  
    use omp_lib  
    real :: a, x(*)  
  
    print*, "id=",omp_get_thread_num()  
  
    !$omp do  
    do i = 1, n  
        x(i)=sqrt(a*i*i)  
    end do  
end subroutine  
  
program main  
    real :: a=9.0, x(16)  
    !$omp parallel  
        call foo(a,x,16)  
    !$omp end parallel  
end program
```

`parallel` creates *team* of threads  
`for / do` workshares implicit tasks among threads

# OpenMP Tasking with Dependences

Makes a bundle of work (task) for asynchronous execution.

- **task** construct -- creates a task out of a code block or function to be queued for execution by any thread of team.
- **depend** clause- orders tasks
  - ❑ between tasks identified by variable
  - ❑ ordering determined by dependence type

```
void init(float *x)      { for(int i=0; i<N; i++) x[i] = i; }  
void axpy(float *x, float *y){ for(int i=0; i<N; i++) y[i] = 2*x[i] + y[i]; }
```

```
int main(){  
    float x[N], y[N];
```

```
{
```

```
    init(y);    // work (task) on y - write to y (out: y)
```

```
    init(x);    // work (task) on x - write to x (out: x)
```

```
    axpy(x,y);  // work (task) on x,y - read x (in: x)  
                // read/write y (inout: y)
```

```
    for(int i=0; i<N; i++){ x[i]=sqrt(y[i]); } // special tasking for loops
```

```
}
```

```
}
```



```
void init(float *x)      { for(int i=0; i<N; i++) x[i] = i; }
void axpy(float *x, float *y){ for(int i=0; i<N; i++) y[i] = 2*x[i] + y[i]; }
```

```
int main(){
    float x[N], y[N];
```

```
#pragma omp parallel master // Get team of threads; 1 thread creates tasks.
```

```
{
```

```
#pragma omp task          // Race here. Any team thread can execute concurrently.
init(y);
```

```
#pragma omp task          // Race here. Any team thread can execute concurrently.
init(x);
```

```
#pragma omp task          // Race here. Any team thread can execute concurrently.
axpy(x,y);
```

```
#pragma omp taskloop // Race here. taskloop chunks iterations into tasks.
```

```
for(int i=0; i<N; i++){ x[i]=sqrt(y[i]); } // Any team thread can execute chunks.
```

```
}
```

```
}
```

# task with depend clause

C/C++

```
void init(float *x)    { for(int i=0; i<N; i++) x[i] = i; }  
void axpy(float *x, float *y){ for(int i=0; i<N; i++) y[i] = 2*x[i] + y[i]; }
```

```
int main(){  
    float x[N], y[N];
```

```
#pragma omp parallel master  
{
```

```
#pragma omp task depend(out: y)    // Task1 writes to y  
    init(y);
```

```
#pragma omp task depend(out: x)    // Task2 writes to x  
    init(x);
```

```
#pragma omp task depend(in: x) depend(inout: y) // Task3 reads x, reads/writes y  
    axpy(x,y);                                // Task3 runs after Task1 & Task2
```

```
#pragma omp taskwait                // Wait for all (sibling) tasks to complete.
```

```
#pragma omp taskloop                // could use depend here and avoid taskwait  
for(int i=0; i<N; i++){ x[i]=sqrt(y[i]); } // execute iteration chunks in parallel
```

```
}  
}
```

Task 1 & 2 can run  
concurrently

# task

program main

integer, parameter :: N=16

real :: x(N), y(N)

call init(x,N) **!! work (task)** on x - **write** to x (**out: x**)

call init(y,N) **!! work (task)** on y - **write** to y (**out: y**)

call axpy(x,y,N) **!! work (task)** on x,y - **read** from x ( **in: x**)  
**!! read/write y (inout: y)**

do i=1,N; x(i)=sqrt(y(i)); enddo **!! special tasking for loops**

end program

```
subroutine init(x,n)
```

```
real :: x(n)
```

```
do i=1,n; x(i) = i; enddo
```

```
end subroutine
```

```
subroutine axpy(x,y,n)
```

```
real :: x(n),y(n)
```

```
do i=1,n; y(i)=2*x(i)+ y(i); enddo
```

```
end subroutine
```

Fortran

OpenMP "speak"

```
program main
```

```
integer, parameter :: N=16
```

```
real                :: x(N), y(N)
```

```
!$omp parallel master    !! Get team of threads; 1 thread creates tasks
```

```
!$omp task               !! Race here. Any team thread can execute concurrently.
```

```
call init(x,N)
```

```
!$omp end task
```

```
!$omp task               !! Race here. Any team thread can execute concurrently.
```

```
call init(y,N)
```

```
!$omp end task
```

```
!$omp task               !! Race here. Any team thread can execute concurrently.
```

```
call axpy(x,y,N)
```

```
!$omp end task
```

```
!$omp taskloop           !! Race here. taskloop chunks iterations into tasks.
```

```
do i=1,N; x(i)=sqrt(y(i)); enddo    !! Any team thread can execute chunks.
```

```
!$omp end parallel master
```

```
end program
```

# task with depend clause

```
program main
```

```
integer, parameter :: N=16
```

```
real :: x(N), y(N)
```

```
!$omp parallel master
```

```
!$omp task depend(out: y)           !! Task1 writes to y.
```

```
call init(x,N)
```

```
!$omp end task
```

```
!$omp task depend(out: x)           !! Task2. writes to x.
```

```
call init(y,N)
```

```
!$omp end task
```

```
!$omp task depend(in: x) depend(inout: y) !! Task3 reads x, reads/writes y.
```

```
call axpy(x,y,N)                   !! Task3 runs after Task1 & Task2.
```

```
!$omp end task
```

```
!$omp taskwait                     !! Wait for all (sibling) tasks to complete.
```

```
!$omp taskloop                     !! could use depend here and avoid taskwait.
```

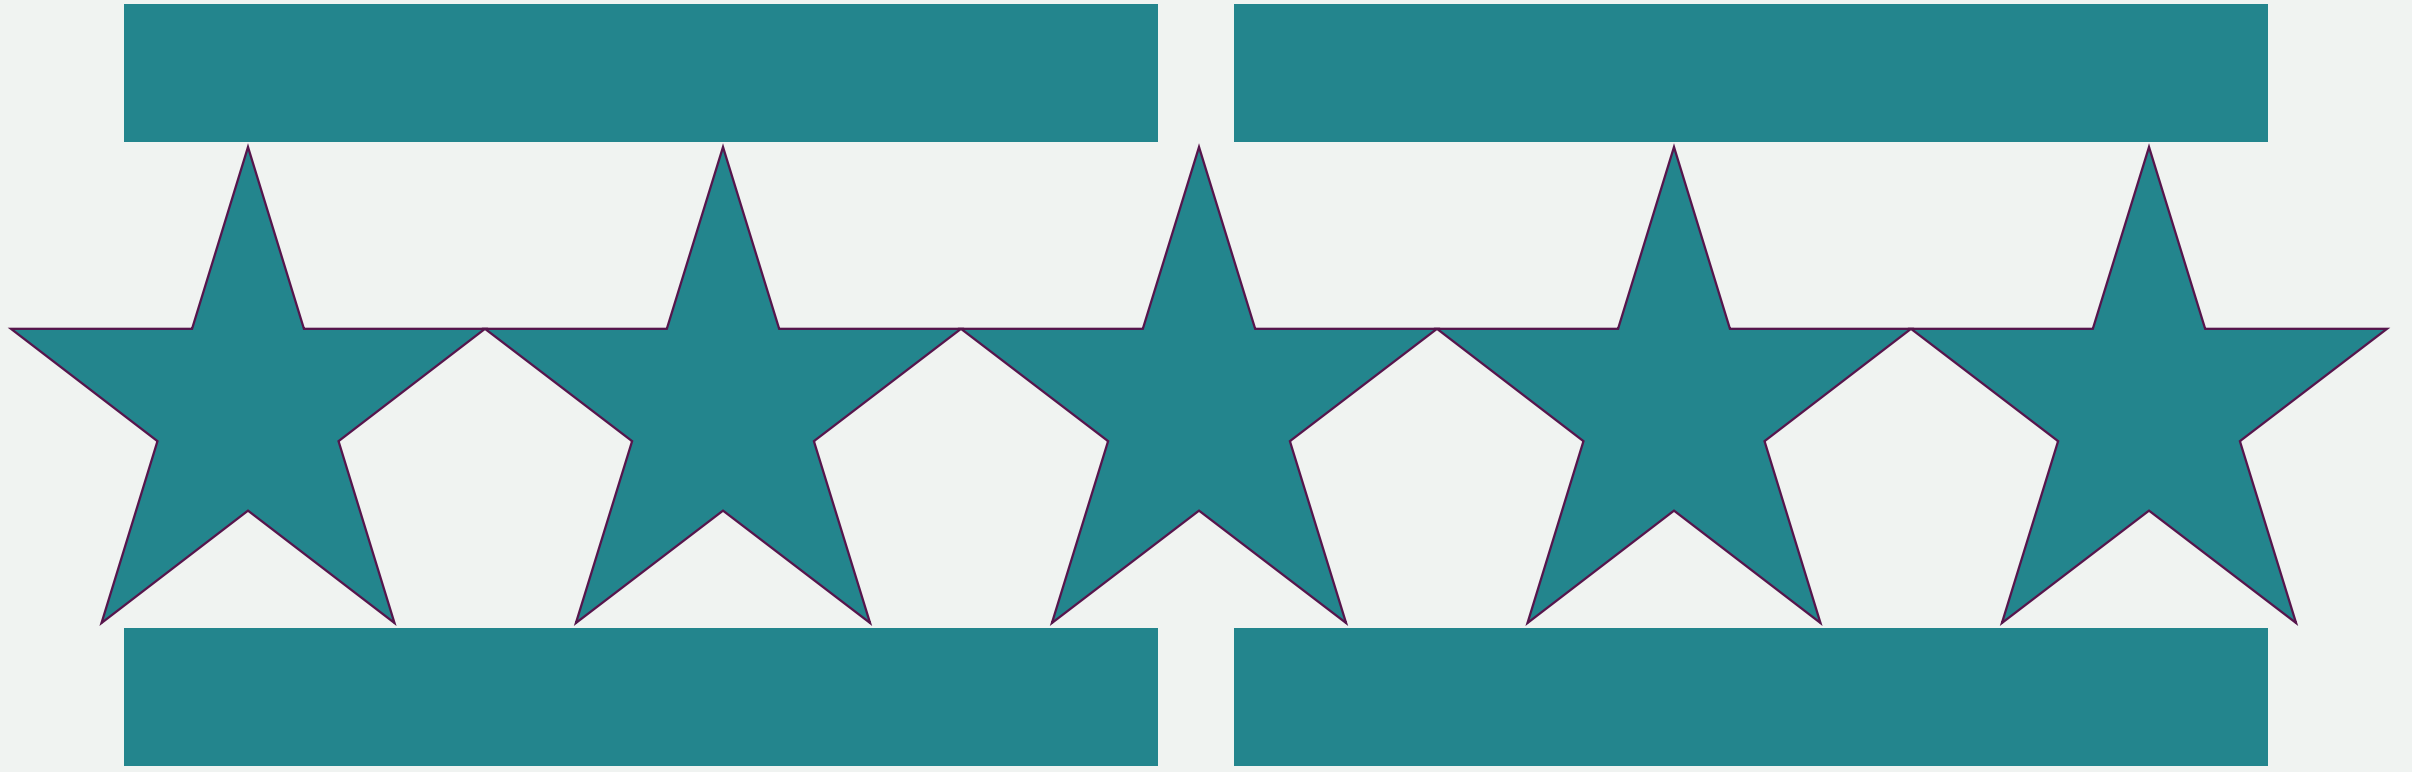
```
do i=1,N; x(i)=sqrt(y(i)); enddo    !! execute iteration chunks in parallel.
```

```
!$omp end parallel master
```

```
end program
```

Task 1 & 2 can run  
concurrently





# OPENMP OFFLOADING



# Agenda

## OpenMP Offloading

- Launch and direct device execution with **target**.
  - Explicit/Implicit **mapping** ("transfer") to/from device.
  - **teams**, **distribute**, and **loop**.
  - Data Persistence: **target data**.
  - Data Movers: **target enter/exit data**, **target update**.
  - **declare target** (ensure function/data execution/access on device).
- Asynchronous Offloading with **nowait** and **depend**.

# target (offloading)

- #pragma omp target

```
int main(){  
    float a=9.0, x[16];  
  
    #pragma omp target  
        foo(a,x);  
  
    #pragma omp target  
    {  
        for(int i=0; i<16; i++)  
            x[i]=sqrt(a*i*i);  
    }  
}
```

**C/C++**

- !\$omp target

```
program main  
    real :: a=9.0, x(16)  
  
    !$omp target  
        call foo(a,x)  
    !$omp end target  
  
    !$omp target  
        do i = 1,16  
            x(i)=sqrt(a*i*i)  
        end do  
    !$omp end target  
  
end program
```

**Fortran**

**Runs on device with only 1 thread.**

Implicit data management: alloc device storage & copy data to device; dealloc & copy from at end of target region.

# target clauses – for beginners

#pragma omp target [ clause [clause ...] ]	C/C++
!\$omp target [ clause [clause ...] ]	Fortran

*device(dev\_id)*

*device #, default 0*

*if(int expr)*

*if false fallback to host (“! 0” in C/C++, logical in F90)*

*map(map-type: var)*  
(Explained in Next Section.)

*alloc/dealloc storage for var on device &  
copy data between host and device.*

on *host var* is called *original* variable/storage  
on *device var* is called *corresponding* variable/storage

*map-type* keywords: *to*, *from*, or *tofrom*

*specify copy direction and implicitly handle storage*

*nowait, depend*

*makes offload async with dependences* (with target task)



# Agenda

## OpenMP Offloading

- Launch and direct device execution with **target**.
- Explicit/Implicit **mapping** ("transfer") to/from device.
  - **teams**, **distribute**, and **loop**.
  - Data Persistence: **target data**.
  - Data Movers: **target enter/exit data**, **target update**.
  - **declare target** (ensure function/data execution/access on device).
- Asynchronous Offloading with **nowait** and **depend**.

# Explicit Data Mapping

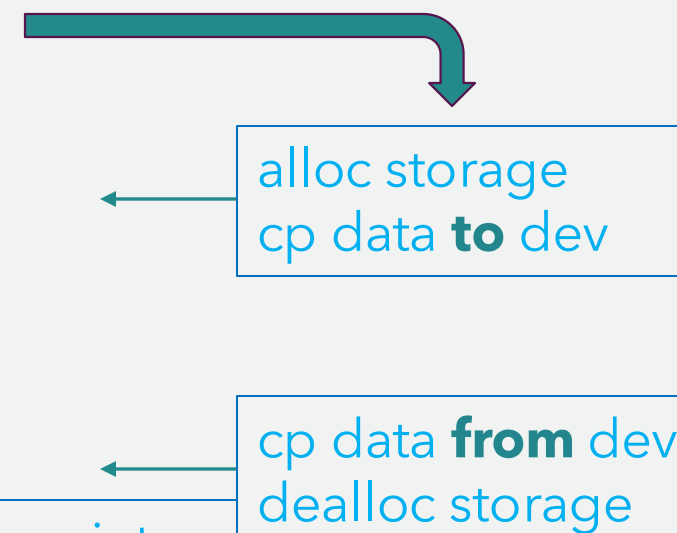
```
#define N 1048576
int main(){
    int scal, array[N], ptee[N], *ptr;

    scal = 0;
    for(int i=0; i<N; i++){ array[i] = 0; ptee[i] = 0; } // inits
    ptr=&ptee[0];

    // Explicit Data Mapping (Defaults to tofrom map-type)

    #pragma omp target map(scal, array, ptr[0:N])
    {
        scal = scal + 1;
        for(int i=0; i<N; i++){ array[i] = array[i] + scal; }
        for(int i=0; i<N; i++){ ptr[i] = ptr[i] + scal; }
    }
}
```

C/C++



Array Section [lb:len:stride] specifies storage pointed to by the pointer.

# Efficient Explicit Data Mapping with map-types

```
#define N 1048576
```

C/C++

```
int main(){
```

```
    int scal, array[N], ptee[N], *ptr;
```

```
    scal = 1;
```

```
    for(int i=0; i<N; i++){ array[i] = 0; ptee[i] = 0; } // inits
```

```
    ptr=&ptee[0];
```

```
// Explicit Data Map(copy scal to, and array and ptee from)
```

```
#pragma omp target map(to: scal) map(from: array, ptr[0:N])
```

```
{
```

```
    for(int i=0; i<N; i++){ array[i] = i + scal; }
```

```
    for(int i=0; i<N; i++){ ptr[i] = i + scal; }
```

```
}
```

```
}
```

# Implicit behavior – first private and data mapping

```
#define N 1048576
int main(){
    int scal, array[N], ptee[N], *ptr;

    scal = 0;
    for(int i=0; i<N; i++){ array[i] = 0; ptee[i] = 0; } // inits
    ptr=&ptee[0];
```

// Implicit Behavior: Works as expected except scalars are firstprivate

```
#pragma omp target:
{
    scal = scal + 1;
    for(int i=0; i<N; i++){ array[i] = array[i] + scal; }
    for(int i=0; i<N; i++){ ptr[i] = ptr[i] + scal; }
}
}
```

C/C++



← *scal* **not**  
copied to host



# Explicit Data Mapping

Fortran

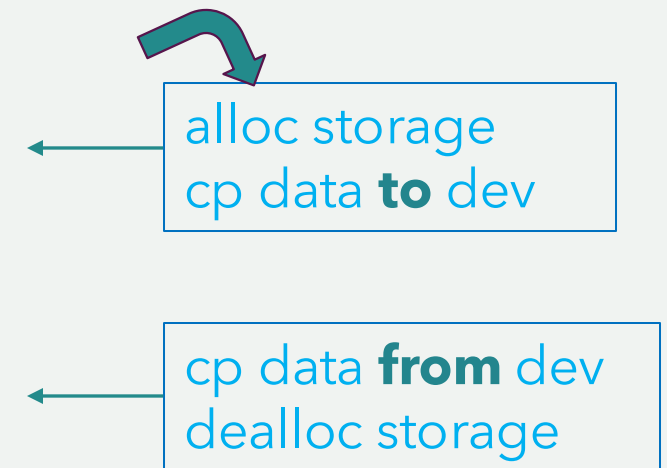
```
program main
  integer, parameter :: N=1048576
  integer          :: scal, array(N)
  integer,pointer   :: ptr(:)
  integer,target    :: ptee(N)

  scal = 0
  do i=1,N; array(i) = 0; ptee(i) = 0; enddo  !! inits
  ptr=>ptee

  !! Explicit Data Mapping (Defaults to tofrom map-type)

  !$omp target map(scal) map(array) map(ptr)
    scal = scal + 1;
    do i=1,N; array(i) = array(i)+1; enddo
    do i=1,N; ptr(i) = ptr(i)+1; enddo
  !$omp end target

end program
```



# Efficient Explicit Data Mapping with map-types

Fortran

```
program main
  integer, parameter :: N=1048576
  integer          :: scal, array(N)
  integer,pointer   :: ptr(:)
  integer,target    :: ptee(N)

  scal = 0
  do i=1,N; array(i) = 0; ptee(i) = 0; enddo  !! inits
  ptr=>ptee

  // Explicit Data Map(copy scal to, and array and ptr from)

  !$ omp target map(to: scal) map(from: array, ptr)

    do i=1,N; array(i) = array(i) + scal; enddo
    do i=1,N; ptr(i)   = ptr(i) + scal; enddo
  !$omp end target

end program
```

# Implicit behavior – first private and data mapping

Fortran

```
program main
  integer, parameter :: N=1048576
  integer          :: scal, array(N)
  integer,pointer   :: ptr(:)
  integer,target    :: ptee(N)

  scal = 0
  do i=1,N; array(i) = 0; ptee(i) = 0; enddo  !! inits
  ptr=>ptee
```

! Implicit Behavior: Works as expected except scalars are firstprivate

!\$omp target

```
  scal = scal + 1;
  do i=1,N; array(i) = array(i)+1; enddo
  do i=1,N; ptr(i) = ptr(i)+1; enddo
```

!\$omp end target

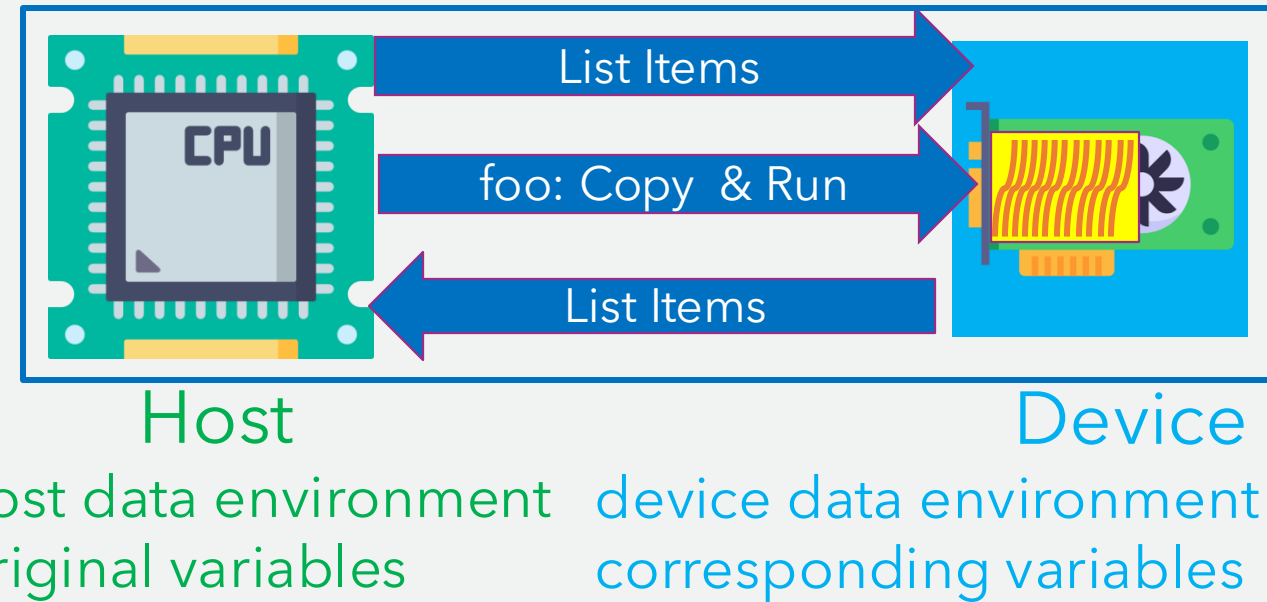
```
end program
```



← *scal* value not  
copied to host

# map clause

What did we just do with map?



`map( action: variables )`

default *action* (map-type) is to and from (tofrom) device

*variables* in lexical extent of target region are implicitly mapped

# map clause syntax

map( *modifier: list items* )

*modifiers*

*list items* (“variables”)

## *map-types*

<b>to</b>	allocate & cp to
<b>from</b>	cp from & deallocate
<b>tofrom</b>	cp to & from
<b>alloc</b>	allocate only ( <b>storage</b> in 6.0)
<i>always</i>	
<b>always</b>	always perform the map-type

## OMP

-----

**scalars**

**arrays**

**pointers**

**objects (structures)**

**array sections**

## C/C++

-----

**int a;**

**int a[8];**

**int \*ptr;**

**struct ...**

**a[start:extent, stride]**

**(used only in OMP directives)**

## F90

-----

**integer :: a**

**integer :: a[N]**

**integer, pointer :: a**

**data type ..**



# Agenda

## OpenMP Offloading

- Launch and direct device execution with **target**.
- Explicit/Implicit **mapping** ("transfer") to/from device.
- **teams, distribute, and loop**.
- Data Persistence: **target data**.
- Data Movers: **target enter/exit data, target update**.
- **declare target** (ensure function/data execution/access on device).
- Asynchronous Offloading with **nowait** and **depend**.



## Parallel **Offloading** – with teams, distribute, and loop

CUDA blocks and threads --> OMP teams and threads

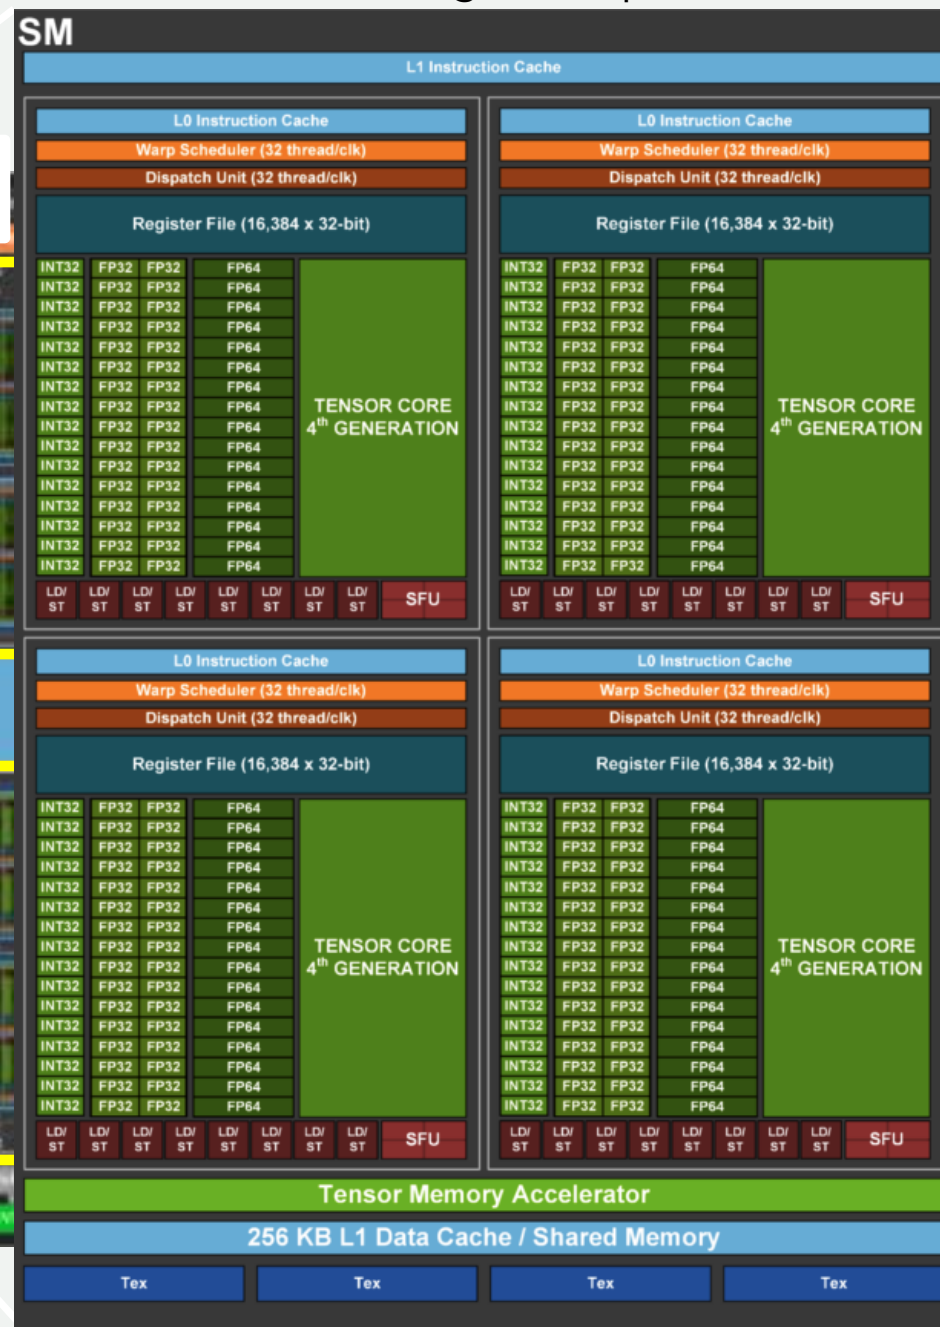
First, GPU architecture:

- ~100 SMs (Streaming Multiprocessors)
- Each SM
  - Is a set of cores (128 FP32, 64 FP64, ..., 4 Tensor)
  - Memory is *shared* within set, but not with other SMs
  - Instruction streams (independent instructions) are scheduled on 32 threads (a warp) to execute simultaneously.

# Hopper Architecture

- 144 SMs (Streaming Multiprocessor).

- Scheduler manages warps on 32 threads



[developer.nvidia.com/blog/nvidia-hopper-architecture-in-depth](https://developer.nvidia.com/blog/nvidia-hopper-architecture-in-depth)



# teams and parallel with target

Creates **teams** for the SMs  
and **threads** within an SM

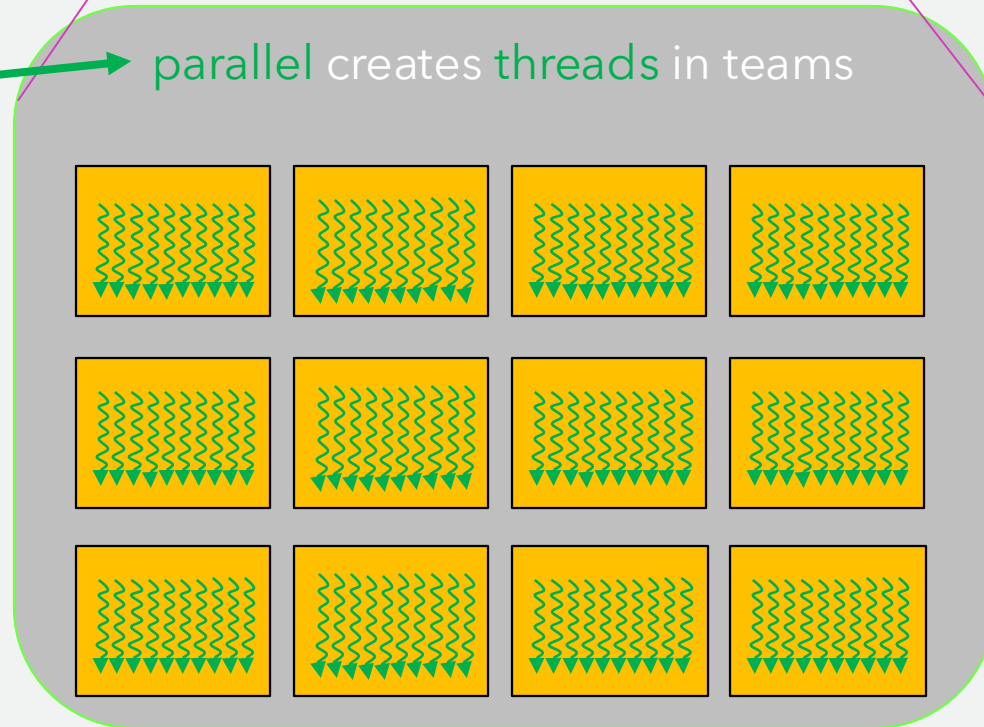
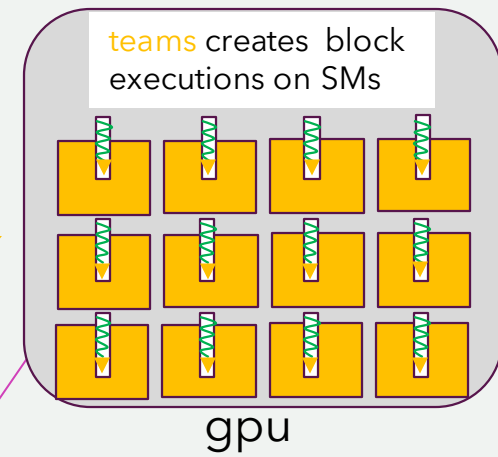
- **teams**: creates teams for execution on SMs
- **parallel**: creates threads within the team
- Default # of threads is implementation defined.

```
#pragma omp target teams  
#pragma omp parallel  
{  
    // code block  
}
```

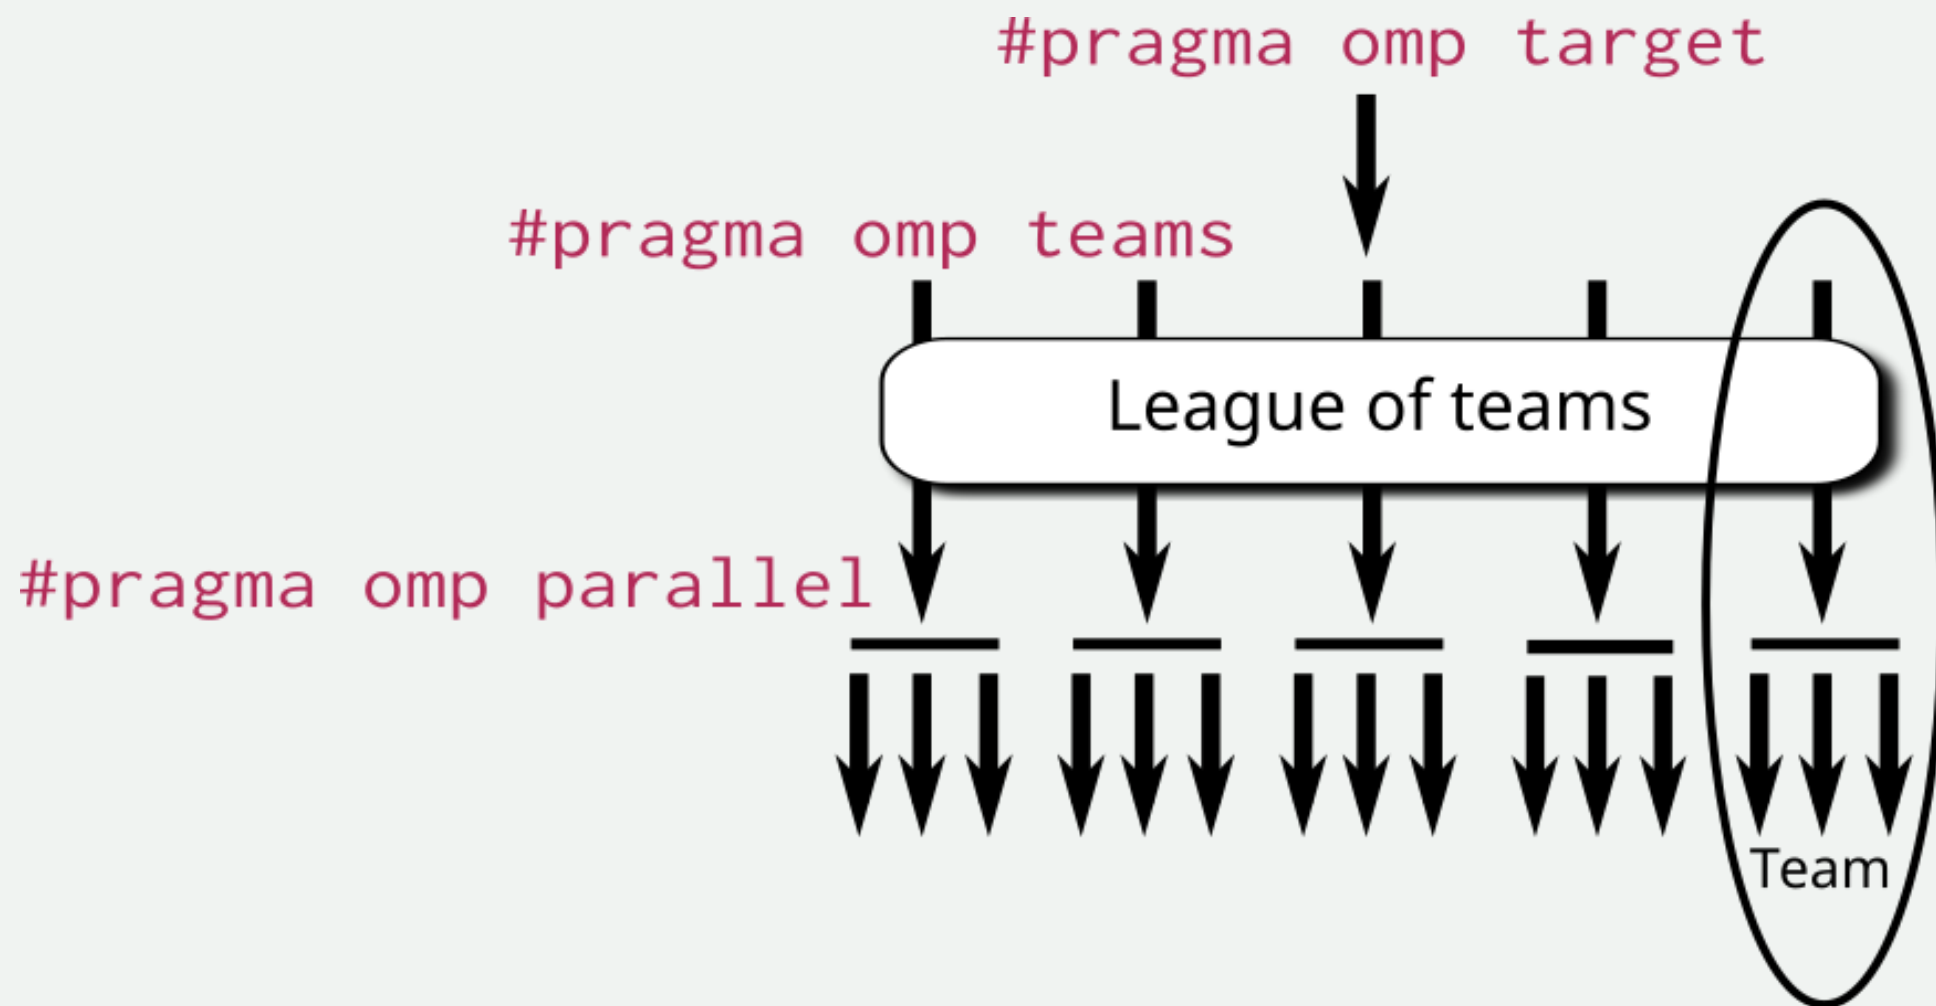
or

```
#pragma omp target teams parallel  
{  
    // code block  
}
```

all executing  
the same  
code block



# League of multi-threaded teams



# work-distribution directives:

partitioning loop iterations across **teams** and **parallel** constructs

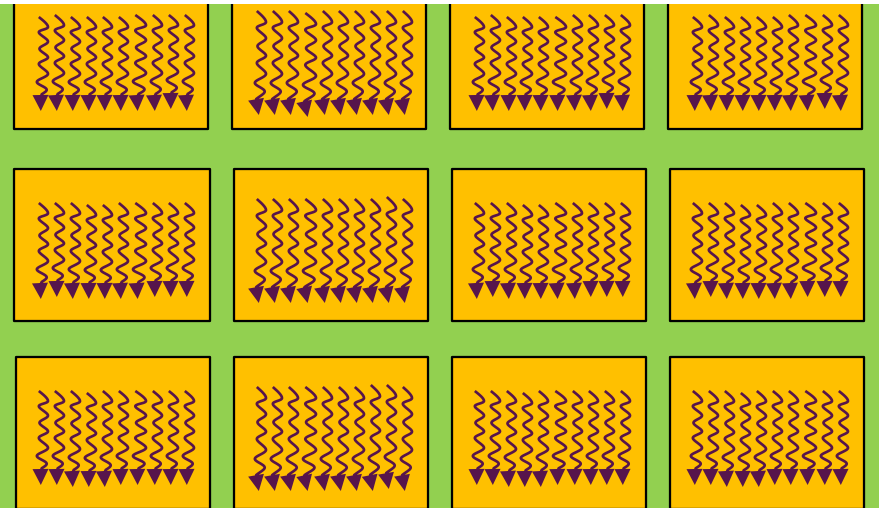
- **distribute** construct specifies iterations of an associated loop are to be distributed across the teams.
- **for** construct is for worksharing across the threads.

```
#pragma omp target teams distribute  
#pragma omp parallel for  
for (...){  
    // code block  
}
```

or

```
#pragma omp target teams distribute parallel for  
for (...){  
    // code block  
}
```

distribute work-sharing across all teams



# Distribute worksharing with target **teams**, **distribute parallel for**, and ("new") **loop**

```
#define N (1 << 27)
int main()
{
    float a=2.0, y[N], x[N];
    for( int i=0; i<N; i++){x[i]=(float)i/N; y[i]=(float)i/N;}
```

```
    #pragma omp target teams distribute parallel for
    for (int i=0; i<N; i++) {y[i]=a*x[i]+y[i];}
```

```
    #pragma omp target teams loop
    for (int i=0; i<N; i++) {y[i]=a*x[i]+y[i];}
```

```
}
```

More control  
with clauses.

More control by  
Implementation.

Number of teams and threads are  
Implementation defined.

# Distribute worksharing with **target teams, distribute parallel for**, and ("new") **loop**

```
program main
  integer, parameter :: N=1024*1024*128
  real :: a=2.0, y(N), x(N), rN=N
  do i=1,N; y(i)=i/rN; x(i)=i/rN; end do

  !$omp target teams distribute parallel do
    do i=1,N; y(i)=a*x(i)+y(i); end do

  !$omp target teams loop
    do i=1,N; y(i)=a*x(i)+y(i); end do

end program
```

More control  
with clauses.

More control by  
Implementation.

Number of teams and threads are  
Implementation defined.



# loop

Allows compiler to control work distribution and worksharing with loop construct.

- CONS: limited user control (clauses)
- PROS: simple syntax, compiler handles the parallelism

```
#pragma omp target teams  
#pragma omp loop  
for (...){  
    // code block  
}
```

or

```
#pragma omp target teams loop  
for (...){  
    // code block  
}
```

```
!$omp target teams  
!$omp loop  
do i=...  
    !! code block  
end do
```

or

```
!$omp target teams loop  
do i=1...  
    !! code block  
end do
```

\* General: can be use in lieu of `parallel for|do` on host.



# Agenda

## OpenMP Offloading

- Launch and direct device execution with **target**.
- Explicit/Implicit **mapping** ("transfer") to/from device.
- **teams**, **distribute**, and **loop**.
- Data Persistence: **target data**.
  - Data Movers: **target enter/exit data**, **target update**.
  - **declare target** (ensure function/data execution/access on device).
- Asynchronous Offloading with **nowait** and **depend**.

# target data – Creating a Data-Persistence Region

```
int main () {  
    int A[2]={10,11};  
  
    #pragma omp target data map(tofrom: A) // map applies to the region {}  
    {  
        // Here: A is allocated & copied to device  
  
        #pragma omp target // (no counterp track)  
        {  
            A[0]++; A[1]++;  
        }  
  
        #pragma omp target // (no counterp track)  
        {  
            A[0]++; A[1]++;  
        }  
  
        // <= Here: On Host A=10,11; On device A=12,13.  
  
        // <= Here: A is copied from device & deallocated, and  
        // Host and Device have same values, A=12,13  
    }  
}
```

C/C++



# target data – Creating a Data-Persistence Region

F O R T R A N

```
program main
  integer :: A(2)=(/10,11/)

  !$omp target data map(tofrom: A) !! map applies to region (here to end target data)
    !! Here: A is allocated & copied to device

    !$omp target no !! (refpointers are not tracked)
      A(1)=A(1)+1; A(2)=A(2)+1
    !$omp end target

    !$omp target no !! (refpointers are not tracked)
      A(1)=A(1)+1; A(2)=A(2)+1
    !$omp end target

    !! <= Here: On Host A=10,11; On device A=12,13.

  !$omp end target data !! <= Here: A is copied from device & deallocated, and
    !! Host and Device have same values, A=12,13
end program
```

# target data Syntax

- map applies to target data region  
employs reference counter

**target data** `map(modifiers: vars)`

## map modifiers

map-type:  
    to, from, tofrom, alloc  
other modifiers: ...

## clauses

device -- device number(s)  
if -- on host | device  
nowait -- async (6.0)  
depend -- task ordering (6.0)  
... -- many others (6.0)

# target enter/exit data -- simple example

```
#include <iostream>
using namespace std;

int main(){
    int A[2]={10,11};

    #pragma omp target enter data map(to: A)

    A[0]=200; A[1]=200;

    #pragma omp target //ref count->no copy
    {A[0]++; A[1]++;}

    #pragma omp target exit data map(from: A)

    cout << A[0]<< " " << A[1] <<endl; //11,12
}
```

program main

C++

integer :: A(2)=(/10,11/)

!\$omp target enter data map(to: A)

A(1)=200; A(2)=200;

!\$omp target !!ref count->no copy

A(1)=A(1)+1; A(2)=A(2)+1

!\$omp end target

!\$omp target exit data map(from: A)

print\*, A(1), A(2) !! 11, 12

end program

# target update -- simple example

```
#include <iostream>
using namespace std;

int main(){
    int A[2]={10,11};

    #pragma omp target data map(A)
    {
        {A[0]=200; A[1]=200;}
        #pragma omp target update to(A)

        #pragma omp target //ref count->no copy
        {A[0]++; A[1]++;}

    }

    cout << A[0]<< " " << A[1] <<endl; //11,12
}
```

C++

program main

F O R T R A N

```
integer :: A(2)=(/10,11/)

!$omp target data map(A)

A(1)=200; A(2)=200;
!$omp target update to(A)

!$omp target !!ref count->no copy
    A(1)=A(1)+1; A(2)=A(2)+1
!$omp end target

!$omp end target data

print*, A(1), A(2) !! 11, 12

end program
```

# Directives for directional data movement: syntax

- map data motion restricted to one direction
- employs reference counter (async with nowait clause)

target\_enter\_data map(type: vars) types: to, alloc

target\_exit\_data map(type: vars) types: from (release, delete)

- Similar to target\_enter/exit\_data, but
- updates without reference counter.

target update to(vars)

target update from(vars)

to and from  
are the data motion clauses

# Compiling code with OpenMP offloading statements

NVHPC <nvc, nvc++, nvfortran>



```
<compiler> -fopenmp -O3 -mp=gpu
```

LLVM <clang, clang++, flang>



```
<compiler> -fopenmp -O3 -fopenmp-targets=nvptx64
```

GNU <gcc, g++, gfortran>



```
<compiler> -fopenmp -O3 -foffload=nvptx-none
```

# LAB I – basic offload

- Create interactive session on gpu compute node.
- Compile and run “hello from gpu” code.
- Compile and run target code block/function for single thread.
- Optimize code for efficient data transfer.
- Compare parallel axpy execution on host and device
- Instrument offloading code to run asynchronously.

```
ssh vista.tacc.utexas.edu
```

```
tar -xvf ~train00/openmp_offload_2025.tar
```

```
cd openmp_offload
```

- Login to Vista
- Untar files and cd to lab directory
- **Read top directory Instructions** file, cd into code directories and **follow instructions in README.md** file.



# Agenda

■ Map Review and Storage



# common ways to change the device data (environment)

Evaluate on your own.

C/C++

```
int scal = 1;
#pragma omp target data map(tofrom: scal) // scal is mapped for region
{
    scal=10;
    #pragma omp target map(always,tofrom: scal) // make device/host consistent (10)
    { scal++; // mapped scal = 11
    } // cp back scal = 11

    scal++;
    #pragma omp target // default: scal is FIRSTPRIVATE, uses Host scal=12
    { scal++; // 13
    } // 12 scal Not Copied back

    scal=100;
    #pragma omp target update to(scal) // consistent scal(100 on host & device)
    scal=500;
    #pragma omp target map(scal) // map ignored
    { scal++; // 101 Increased 100 to 101 on Device
    } // 500, Not Copied back
} // 101 Copied back, same on Host and Device
```



# common ways to change the device data (environment)

Evaluate on your own.

F O R T R A N



```
integer :: scal = 1
```

```
!$omp target data map(tofrom: scal)    !! scal is mapped for region
```

```
scal=10
```

```
!$omp target map(always,tofrom: scal) !! make device/host consistent (10)
```

```
    scal=scal+1                !! mapped scal = 11
```

```
!$omp end target
```

```
scal=scal+1
```

```
!$omp target    !! default: scal is FIRSTPRIVATE, uses Host scal=12
```

```
    scal=scal+1    !! 13
```

```
!$omp end target  !! 12 scal Not Copied back
```

```
scal=100
```

```
!$omp target update to(scal)  !! consistent scal(100 on host & device)
```

```
scal=500
```

```
!$omp target map(scal)        !! map ignored
```

```
    scal=scal+1                !! 101 Increased 100 to 101 on Device
```

```
!$omp end target              !! 500, Not Copied back
```

```
!$omp end target data          !! 101 Copied back, same on Host and Dev
```

# device-only storage -- like cuda, uses device pointer

```
int main(){
    double Val, *d_B, *d_C;  // (d_ = device pointers)

    int dev_id = omp_get_default_device();    // device id number
    d_B=(double *)omp_target_alloc(N*N*sizeof(double),dev_id);
    d_C=(double *)omp_target_alloc(N*N*sizeof(double),dev_id);

    #pragma omp target map(from: Val) is_device_ptr(dev_B,dev_C) device(dev_id)
    {
        //specify dev pointers & their device

        for(int i=0; i<N; ++i)  // B,C storage used (no data transfer)
            for(int j=0; j<N; ++j)
                {d_B[i*N+j]=i*j; d_C[i*N+j]=i*j+1;}

        ...
        Val=3.14;
    }

    omp_target_free(dev_B, dev_id);
    omp_target_free(dev_C, dev_id);
}
```

C/C++

## alloc maptype for temporary space

```
float x[N], y[N];C/C++  
...  
#pragma omp target ... map(alloc: x) map(tofrom: y)  
{  
    for(int i=0;i<n;i++) x[i]=sin(i%4)*cos(i%6);  
  
    for(int i=0;i<n;i++) y[i]=a*x[i]+y[i];  
}
```



# Agenda

## OpenMP Offloading

- Launch and direct device execution with **target**.
- Explicit/Implicit **mapping** ("transfer") to/from device.
- **teams**, **distribute**, and **loop**.
- Data Persistence: **target data**.

Data Movers: **target enter/exit data**, **target update**.

■ **declare target** (ensure function/data execution/access on device).

Asynchronous Offloading with **nowait** and **depend**.

# declare target – declare functions as device procedures

- Procedures in the same “compile” unit are automatically compiled for devices.
- “Externally” compiled procedures must be **declared** for **target** devices.

```
void fun();  
  
int main(){  
    #pragma omp target  
    fun();  
}
```

C/C++

```
#include <omp.h>           // Difference Compile Unit  
#include <stdio.h>  
  
void fun(){  
    printf("%d\n", omp_is_initial_device()); // 0 on GPU, 1 on HOST  
}  
  
#pragma omp declare target (fun) // Compile for GPU, too.
```

C/C++

# declare target -- Multiple ways to do it (global vars, too)

```
#pragma omp begin declare target
```

C/C++

```
static float a_global=2.0f;  
void foo(){printf("external compiled unit\n");}
```

This is for declaring everything within a region

```
#pragma omp end declare target
```

```
static float a_global=2.0f;  
void foo(){printf("external compiled unit\n");}
```

C/C++

```
#pragma omp declare target (a_global, foo)
```

This is for declaring specific vars and funcs.

```
static float a_global=2.0f;  
void foo(){printf("external compiled unit\n");}
```

C/C++

```
#pragma omp declare target to(a_global, foo)
```

Same as above, using clause (**to** is default).

# declare target – declare functions as device procedures

- Procedures in the same “compile” unit are automatically compiled for devices.
- “Externally” compiled procedures must be **declared** for **target** devices.

```
program main
  use functions
  !$omp target
    call fun()
  !$omp end target
end program
```

F

```
module functions      !! Different “Compile” unit.
  use omp_lib
  contains
  subroutine fun()
    !$omp declare target (fun)    !! Compile for GPU, too.
    print*, omp_is_initial_device() !! F on GPU, T on HOST
  end subroutine
end module
```

F



# declare target -- Fortran

```
module functions
contains
  subroutine fun()
    !$omp declare target to(fun)
    print*, "externally compiled"
  end subroutine
end module
```

F 90

```
program main
  !$omp target
  call fun()
!end target
end program
```

```
module functions
integer,parameter :: N=8
real              :: v(N)
contains
  subroutine fun(one)
    integer, intent(in) :: one
    !$omp declare target enter(fun,v)
    print*, "externally compiled"
    v(1:N)=one
  end subroutine
end module
```

F 90

Include immediately  
after specifications.

List function and  
global variables in  
to/enter clause.

```
program main
  use functions
  integer :: one=1
  !$omp target
  call fun(one)
!$omp end target
!$omp target update from(v)
  print*,v
end program
```

# declare target syntax

#pragma omp **begin** declare target

//procedures & static variables here

#pragma omp **end** declare target

C/C++ only  
6.? Fortran

#pragma omp declare target (*ext\_list*)

Use !\$omp for F90

*ext\_list*: list of procedures(pcr) and/or static variables(var)

#pragma omp declare target *clauses*

Use !\$omp for F90

*clauses*:      **to/enter**(*ext\_list*)    avail. entire prog  
                 link(    *var\_list*)        avail. when mapped  
                 device(    *expr*)    device specific (**id**)  
                 indirect(*prc\_list*)    for func. pointers

**to**(<v5.2)    **enter**(>=v5.2)

# Asynchronous Offloading

```
ndev=omp_get_num_devices(); // Number of devices, ids=0,...,ndev-1
n = N/ndev;
```

C/C++

```
#pragma omp parallel num_threads(ndev)
#pragma master
{
    for(int id=0; id<ndev; id++)
    {
        //simple,unblanced partitioning
        ib = id*n; ie = ib+n;
        if(id==ndev-1){ ie=N; n=N-ib;} //final iter.
        #pragma omp target teams distribute parallel for \
            nowait device(id) map(y[ib:n],x[ib:n]) //ib=offset,n=sz
        for(int i=ib; i< ie; i++) y[i] = a * x[i] + y[i];
    }
    // parallel barrier waits for tasks to complete
}
```

Could do this without *parallel master*, but use a *taskwait* at end to synchronize.

# Asynchronous Offloading

```
ndev=omp_get_num_devices()    !! Number of devices, ids=0,...,ndev-1
n = NT/ndev
```

```
!$omp parallel num_threads(ndev)
```

```
!$omp master
```

```
do id=0,ndev-1
```

```
!!partitioning: simple,unblanced
```

```
ib = 1 + id*n; ie = ib + n-1
```

```
if(id == ndev-1) then; ie=NT; endif    !!final iter.
```

```
!$omp target teams distribute parallel for &
```

```
!$omp&    nowait device(id) map(y(ib:ie),x(ib:ie))
```

```
do i=ib,ie; y(i) = a * x(i) + y(i); enddo
```

```
enddo
```

```
!$omp end master
```

```
!$omp end parallel    !! parallel barrier waits for tasks to complete
```

F

**begin**  
**end**  
Fortran  
Array  
Syntax

Could do this without parallel master, but use a taskwait at end to synchronize.

# Extras



# Map: Reference count

- On entry to device environment:

Atomic Operation

- If a corresponding list item is not present in the device data environment, then:
  - A new list item corresponding to original list item (on host) is created in the device data environment;
  - The corresponding list item has a reference count that is initialized to zero; and
  - The value of the corresponding list item is undefined;
- If ref count is not incremented due to map clause, it is incremented by 1

- On exit from device environment:

Atomic Operation

- if map-type is **delete** ref count is set to 0
- if map-type is not **delete** the ref count is decremented by 1 (min 0)
- If the reference count is zero then the corresponding list item is removed from the device data environment.

# Direct Device Allocation: omp\_target\_alloc() + is\_device\_ptr

```
#include <omp.h>
int main(){
int a=2, N=1<<4; // 16
int dev_no = omp_get_default_device();

int *y =(int*)      malloc(N*sizeof(N) );
int *x_d=(int*)omp_target_alloc(N*sizeof(N), dev_no);

for(int i=0;i<N;i++){ y[i]=1; }

#pragma omp target is_device_ptr(x_d)
for(int i=0;i<N;i++) x_d[i]=1;      //INIT

#pragma omp target is_device_ptr(x_d) \
        map(tofrom: y[0:N])
for(int i=0;i<N;i++) y[i]=a*x_d[i]+y[i]; //AXPY

omp_target_free(x_d, dev_no);
}
```

C/C++

# omp\_target\_alloc syntax

Syntax: `void* omp_target_alloc(size_t size, int device_num);`

`returns` the device address of  
a storage location of `size bytes`

`device_num` : less than the result of  
`omp_get_num_devices()` or the result of a call to  
`omp_get_initial_device()`.

```
int init_dev = omp_get_initial_device();
int ndevs = omp_get_num_devices();
cout << "init_dev# = " << init_dev << " # non-host_devs" << ndevs << endl;
```

OUTPUT: init\_dev# = 3 # non-host\_devs3



# Direct memory copy with omp\_target\_memcpy

C/C++

```
void get_dev_cos(double *mem, int s){
    int h, t, i;
    double * mem_dev_cpy;
    h = omp_get_initial_device();
    t = omp_get_default_device();

    mem_dev_cpy = (double *)omp_target_alloc( sizeof(double) * s, t);

    /* dst src */
    omp_target_memcpy(mem_dev_cpy, mem, sizeof(double)*s,
                      0, 0,
                      t, h);

    #pragma omp target is_device_ptr(mem_dev_cpy) device(t)
    #pragma omp teams distribute parallel for
    for(i=0;i<s;i++){ mem_dev_cpy[i] = cos((double)i); } /* init data */

    /* dst src */
    omp_target_memcpy(mem, mem_dev_cpy, sizeof(double)*s,
                      0, 0,
                      h, t);
    omp_target_free(mem_dev_cpy, t);
}
```

# Verifying OpenMP Offloading

- When the (generic) OpenMP flag is on (e.g. usually `-fopenmp`) and a compiler doesn't understand the target compiler flag it **defaults to fallback execution of the region on the host.**

How to check if target is offloading to GPU.

```
#include <stdio.h>
#include <omp.h>
int main(){
    #pragma omp target
    {printf("%d\n", omp_is_initial_device());} // 0 run on GPU, 1 run on HOST
}
```

C/C++

```
program main
    use omp_lib
    !$omp target
        print*, omp_is_initial_device()      !! F run on GPU, T run on HOST
    !$omp end target
end program
```

Fortran

# Verifying OpenMP Offloading

Useful offloading test.

```
#include <stdio.h>
#include <omp.h>
int main(){
    #pragma omp target
    {printf("ranOnGPU=%c T=true\n", (omp_is_initial_device()==0) ? 'T':'F');}
}
```

**C/C++**

```
program main
    use omp_lib
    !$omp target
        print*, "ranOnGPU= ", .not. omp_is_initial_device(), " T=TRUE, F=FALSE"
    !$omp end target
end program
```

**Fortran**